

Burrows Wheeler Transform

Janosch Ruff

Georg-August-Universität Göttingen

January 12, 2021

Structure

- 1 Compression
- 2 Burrows Wheeler Transform
- 3 Inverting the Burrows-Wheeler Transform
- 4 Outlook

Core idea

Definition

***Data compression** is the process of encoding information by using fewer bits than its original representation.*

- Idea: Transform a string to make it "more compressible".
- Question: Why is a string more compressible than another?

What is compressible?

Kolmogorov Complexity

- "Kolmogorov complexity is a notion every computer scientist should know about but not every computer scientist knows."

Definition

*Fix a programming Language. Then the **Kolmogorov Complexity** $C(\cdot)$ for a string ω is defined by*

$$C(\omega) := \text{Bit length of the shortest program that outputs } \omega \quad (1)$$

- "Kolmogorov complexity represents the limit for optimal compressors and thus is the Gold Standard for compression."

Compression examples

Stochastic Process



Figure 1: Flipping a coin

Example 1 a simple pattern: 10101010..101010

Example 2 the first 128 bits of π :

1100100100001111110110101010001000100001011010001100001000110
1001100010011000110011000101000101110000000110111000001110011
010001..

Motivation

- Observation: A string with low Kolmogorov Complexity has a "structure" but some structures are simpler to see than others.
- Idea: Construct an algorithm to transform a string that discovers the patterns given by "low" Kolmogorov Complexity.
- Aim: Exploit discovered patterns to improve compression.

Kolmogorov Complexity

Properties

- The Kolmogorov Complexity is not computable.
- The Kolmogorov Complexity of a string is at most the length of the string itself $C(\omega) \leq |\omega| + \mathcal{O}(1)$ (ω hard coded).
- No "magic universal compressor" exists and $\exists \omega$ s.t. $C(\omega) \geq |\omega|$ (There are always strings that are not compressible).
- Let ϕ be a computable bijection then

$$C(\phi(\omega)) = C(\omega) + C(\phi) + \mathcal{O}(1) = C(\omega) + \mathcal{O}(1)$$
- An incompressible string ξ with $C(\xi) = |\xi| + \mathcal{O}(1)$ is equivalent to an algorithmic random string.

Sorting

A simple idea

- Idea: Enforce a structure by sorting the string.
- Example original string:
"kolmogorov complexity is a notion every computer scientist should know about but not every computer scientist knows."
- Example string after sorting:
"aabbccccdeeeeeeeeghiiiiikkllllmmmmnnnnnnnn
oooooooooooooooooppprrrrrrsssssstttttttttttuuuuuvvvwwwxyyy"
- Result: Long runs of letters and thus easier to compress.
- Problem: How to recover the original string from the string after being sorted?

Sorting

A little bit more sophisticated

- Next idea: Sort the word to discover potential patterns.
- Implementation: Output the letter which occurs in the original string before the sorted letter.
- Example: "BANANA\$" \xrightarrow{BWT} "ANNB\$AA"

\$BANANA

A\$BANAN

ANA\$BAN

ANANA\$B

BANANA\$

NA\$BAN

NANA\$BA

Properties

Questions

- Can the sorting be done in linear time $\mathcal{O}(n)$?
- Is it easy to find a good compressor φ for $BWT(\omega)$?
- Is the transformed string invertible $BWT^{-1}(BWT(\omega)) = \omega$?

Sorting

In linear time

- Notice that the green colored substrings is a Suffix Tree/
Suffix Array of the string "BANANA\$"

\$BANANA

A\$BANAN

ANA\$BAN

ANANA\$B

BANANA\$

NA\$BANA

NANA\$BA

- Lemma: We can construct a Suffix Array in linear time.

Transformation

Using Suffix Array

- Concatenation $\alpha = \omega \cdot \$ \cdot \omega$
(e.g. $\omega = \text{"BANANA"} \xrightarrow{\alpha} \text{"BANANA\$BANANA"}$)
- Create a suffix array of $\omega \cdot \$$
(e.g. $\omega = \text{"BANANA"} \xrightarrow{SA} [6,5,3,1,0,4,2]$)

[6] \$BANANA

[5] A\$BANAN

[3] ANA\$BAN

[1] ANANA\$B

[0] BANANA\$

[4] NA\$BANA

[2] NANA\$BA
- $\eta[i] = \alpha[SA[i] + n]$ where $|\omega| = n$ and $0 \leq i \leq n$.

Transformation

Algorithm

Algorithm 1 BWT(ω)

```
1:  $\alpha \leftarrow \omega\$ \omega$ 
2:  $n \leftarrow |\omega|$ 
3:  $\Gamma \leftarrow \text{SuffixArray}(\omega\$)$ 
4:  $\eta \leftarrow \epsilon$ 
5:  $i \leftarrow 0$ 
6: while  $i \leq n$  do
7:    $\eta[i] \leftarrow \alpha[\Gamma[i] + n]$ 
8:    $i \leftarrow i + 1$ 
9: end while
10: return  $\eta$ 
```

Properties

Answers

- Can the sorting be done in linear time $\mathcal{O}(n)$? ✓
- Is it easy to find a good compressor φ for $BWT(\omega)$?
- Is the transformed string invertible $BWT^{-1}(BWT(\omega)) = \omega$?

Compressibility

Why $BWT(\omega)$ is easier to compress?

- Observation: $\eta[i]$ is the letter which occurs before $\omega_{sort}[i]$.
- This shows that η depends on a sorted list
(e.g. consider the word "banana", then "n" occurs always before "a" and thus we can expect *runs* of "n".)
- There exists no "magic universal compressor" but we can apply a custom-made compression scheme φ exploiting the properties of a transformed string with predictable properties.
- Kolmogorov Complexity tells us $C(BWT(\omega)) = C(\omega) + \mathcal{O}(1)$ but BWT clusters the letters based on the structure of ω and reveals hidden patterns what makes it easier to compress for φ .

Properties

Answers

- Can the sorting be done in linear time $\mathcal{O}(n)$? ✓
- Is it easy to find a good compressor φ for $BWT(\omega)$? ✓
- Is the transformed string invertible $BWT^{-1}(BWT(\omega)) = \omega$?

Inverse Functions

Basics




Figure 2: Example of an inverse function f

Unsort the sorted string

Available information

- Observation: $\eta[i]$ is the letter which occurs before $\omega_{sort}[i]$.
- This entails that $\eta[0]$ is the letter which occurs before $\omega_{sort}[0] = \$$ hence $\eta[0] = \omega[n-1]$ the last letter of ω
\$BANANA
A\$BANAN
ANASBAN
ANANASB
BANANAS
NASBANA
NANASBA
- Idea: Reconstructing ω from η by finding the letter occurring before the letter $\eta[i]$ in ω .

First Last Column

Indexing

- Observation: We know in which row of the sorted list (first column) the element $\eta[k]$ can be found (because it is sorted).
- The last element in this row indicates the letter which occurs before $\omega_{sort}[k]$

F..L

\$₁..A₁

A₁..N₁

A₂..N₂

A₃..B₁

B₁..\$₁

N₁..A₂

N₂..A₃

- Observation: The structure of the first row is completely predictable since it is sorted and thus only the last column is needed to recreate ω from η hence, $\exists BWT^{-1}(\eta) = \omega$.

BWT⁻¹

Algorithm

Algorithm 2 BWT⁻¹(η)

```
1:  $\pi \leftarrow \text{permutation}(\eta)$ 
2:  $\omega \leftarrow \epsilon$ 
3:  $n \leftarrow |\eta|$ 
4:  $i \leftarrow n - 1$ 
5:  $k \leftarrow 0$ 
6: while  $i \geq 0$  do
7:    $\omega[i] \leftarrow \eta[k]$ 
8:    $i \leftarrow i - 1$ 
9:    $k \leftarrow \pi[k]$ 
10: end while
11: return  $\omega$ 
```

Calculating the permutation

Algorithm

Algorithm 3 permutation(η)

```
1:  $\pi \leftarrow [0] \cdot |\eta|$ 
2:  $\Lambda, \vartheta \leftarrow [0] \cdot |\Sigma|$ 
3: for  $i = 0, 1, \dots, |\eta| - 1$  do
4:    $\Lambda[\eta[i]] \leftarrow \Lambda[\eta[i]] + 1$ 
5: end for
6: for  $i = 1, 2, \dots, |\eta| - 1$  do
7:    $\vartheta[i] \leftarrow \vartheta[i - 1] + \Lambda[i - 1]$ 
8: end for
9: for  $i = 0, 1, 2, \dots, |\eta| - 1$  do
10:   $\pi[i] \leftarrow \vartheta[\eta[i]]$ 
11:   $\vartheta[\eta[i]] \leftarrow \vartheta[\eta[i]] + 1$ 
12: end for
13: return  $\pi$ 
```

Properties

Answers

- Can the sorting be done in linear time $\mathcal{O}(n)$? ✓
- Is it easy to find a good compressor φ for $BWT(\omega)$? ✓
- Is the transformed string invertible
 $BWT^{-1}(BWT(\omega)) = \omega$? ✓

Further upgrades

What can be improved?

- Question: Is it possible to do the BWT without an additional special character \$?




Figure 3: Example sketch of a bijective function

Bijection

Lyndon Words

Definition

*A **Lyndon Word** is strictly smaller than any of its proper suffixes and strictly smaller in lexicographic order than all of its rotations.*

- Lyndon Words of the binary alphabet:
0, 1, ~~00~~, 01, ~~10~~, ~~11~~, ~~000~~, 001, ~~010~~, 011, ~~100~~, ~~101~~, ~~110~~, ~~111~~, ~~0000~~, 0001, ..
- Chen–Fox–Lyndon Theorem: Every word can be uniquely factored into $\omega_1\omega_2..\omega_m$ Lyndon words where $\omega_i \geq_{\text{lex}} \omega_{i+1}$.
- "BANANA" $\xrightarrow{\text{Factors}}$ "B" $>_{\text{lex}}$ "AN" \geq_{lex} "AN" $>_{\text{lex}}$ "A"
- Compute a list of all Lyndon factors of ω and their respective rotations and sort this list by $u <_{\omega} v$ such that $uuu.. <_{\text{lex}} vvv..$
- $B|AN|AN|A \xrightarrow{\text{Rotations}} B|AN, NA|AN, NA|A$
 $\xrightarrow{\text{Sort}} A <_{\omega} AN \leq_{\omega} AN <_{\omega} B <_{\omega} NA \leq_{\omega} NA$
 $\xrightarrow{\text{Chars}} ANNBAA$

BBWT

Algorithm

Algorithm 4 BBWT(ω)

```
1:  $\omega_1\omega_2..\omega_m \leftarrow \text{LyndonFactors}(\omega)$ 
2:  $\chi \leftarrow \emptyset$ 
3: for  $i = 1, 2, \dots, m$  do
4:    $\chi \cup \text{CyclicRotations}(\omega_i)$ 
5: end for
6:  $\tau \leftarrow \text{OmegaSort}(\chi)$ 
7:  $\eta \leftarrow \text{LastChars}(\tau)$ 
8: return  $\eta$ 
```

- Open Problem: Can line 6 the sorting based on comparing infinite periodic repetitions be done in linear time?

Inverting BBWT

Key Observations

- Observation: If $\eta = BBWT(\omega)$, then the last character of ω_m (and hence of ω) is $\eta[0]$.
- We can reconstruct ω by reconstructing the i -th Lyndon Factor ω_i by using the computation of the Algorithm permutation.
- Question: How do we notice that we completed the i -th Lyndon Factor and where do we proceed?
- Idea: Keep track of visited letters in η telling us if we would "recycle" this letter and we have to skip it in which case we notice that the last letter of ω_{i-1} is next to the last letter of ω_i .

BBWT⁻¹

Algorithm

Algorithm 5 BWT⁻¹(η)

```
1:  $\pi \leftarrow \text{permutation}(\eta)$ 
2:  $\nu \leftarrow [\mathbf{False}] \cdot |\eta|$ 
3:  $j, k \leftarrow 0$ 
4: for  $i = |\eta| - 1, |\eta| - 2, \dots, 0$  do
5:   while  $\nu[k] \neq \mathbf{False}$  do
6:      $j \leftarrow j + 1$ 
7:      $k \leftarrow j$ 
8:   end while
9:    $\omega[i] \leftarrow \eta[k]$ 
10:   $\nu[k] \leftarrow \mathbf{True}$ 
11:   $k \leftarrow \pi[k]$ 
12: end for
13: return  $\omega$ 
```

Conclusion

Remarks

File	Size	BWT-compression		S-compression		Gain	
		bytes	ratio	bytes	ratio	absolute	relative
BIB	111,261	32,022	28.78%	31,197	28.04%	0.74%	2.58%
BOOK1	768,771	242,857	31.59%	235,913	30.69%	0.90%	2.86%
BOOK2	610,856	170,783	27.96%	166,881	27.32%	0.64%	2.28%
GEO	102,400	66,370	64.81%	66,932	65.36%	-0.55%	-0.85%
NEWS	377,109	135,444	35.92%	131,944	34.99%	0.93%	2.58%
OBJ1	21,504	12,727	59.18%	12,640	58.78%	0.40%	0.68%
OBJ2	246,814	98,395	39.87%	94,565	38.31%	1.55%	3.89%
PAPER1	53,161	19,816	37.28%	18,931	35.61%	1.66%	4.47%
PAPER2	82,199	28,084	34.17%	27,242	33.14%	1.02%	3.00%
PAPER3	46,526	18,124	38.95%	17,511	37.64%	1.32%	3.38%
PAPER4	13,286	6,047	45.51%	5,920	44.56%	0.96%	2.10%
PAPER5	11,954	5,815	48.64%	5,670	47.43%	1.21%	2.49%
PAPER6	38,105	14,786	38.80%	14,282	37.48%	1.32%	3.41%
PIC	513,216	59,131	11.52%	52,406	10.21%	1.31%	11.37%
PROGC	39,611	15,320	38.68%	14,774	37.30%	1.38%	3.56%
PROGL	71,646	18,101	25.26%	17,916	25.01%	0.26%	1.02%
PROGP	49,379	13,336	27.01%	13,010	26.35%	0.66%	2.44%
TRANS	93,695	22,864	24.40%	22,356	23.86%	0.54%	2.22%
Total	3,251,493	980,022	30.14%	950,090	29.22%	0.92%	3.05%
Median	76,923	21,340	36.60%	20,644	35.30%	0.94%	2.58%

Figure 4: Comparing performance between BWT-based compression and S-based (BBWT) compression of the Calgary corpus.

Sources

Further material

- The original paper on Burrows-Wheeler Transform "[A Block-sorting Lossless Data Compression Algorithm](#)" by [Burrows & Wheeler](#) was published in 1994.
- The paper "[A Bijective String Sorting Transform](#)" by [Gil & Scott](#) has been used as the main source for the presentation and might serve for further studies in bijective BWT.
- The data compressor [bzip2](#) uses BWT in combination with [Move-to-front transform](#) and [Huffman coding](#) to compress files.
- Another home for BWT has been found in the field of Bioinformatics in the form of FM-index [here smoothly introduced by Ben Langmead](#) to approximate string alignment.
- As a light introduction into the world of Kolmogorov Complexity or rather Algorithmic Information Theory the [lecture given by Shai Ben-David](#) is a nice starting point.

Appendix

Duval's Algorithm

- The algorithm originally published in the paper "**Factorizing words over an ordered alphabet**" by Duval in 1983 computes the Lyndon Factors of the word ω in linear time $\mathcal{O}(|\omega|)$.
- A complete python implementation of the two transform versions and further compression techniques is provided [here](#).

#Python implementation of Duval's Algorithm for Lyndon Factors

```
def duval(w):  
    i = 0  
    n = len(w)  
    factors = []  
    while(i < n):  
        j = i + 1  
        k = i  
        while(j < n and w[k] <= w[j]):  
            if(w[k] < w[j]):  
                k = i  
            else:  
                k += 1  
            j += 1  
        while(i <= k):  
            factors.append(w[i:i+j-k])  
            i += j-k  
    return(factors)
```